


d=7, r=7, p=0.004 — parameter-matched Mamba vs GRU, 10M training shots, Adam 3e-3, batch 10,000

Same data, same input, matched parameters:
only the recurrent layer differs



Train and validation coincide (0.1031 vs 0.1037):
the model underfits, it does not overfit

