



Experiment 8 — RCNN teacher at $r=5$

Surface code $d=5$, $p=0.01$, $r=5$ · 200,000-shot validation partition [20.0M, 20.2M) · 3 seeds, bar = full range, marks = individual seeds



p_L is from each run's re-scored `.best checkpoint (val_scores_best_ckpt.csv)`, not the weights left in memory when `fit()` returned — the two agree to 6 d.p. for all three seeds here.
RCNN 69,485 parameters. Only the 10M rung has been run so far; the lower rungs and 20M are pending. MWPM is recomputed on this validation partition — the 0.08910 printed in the training logs belongs to the old $r=5$ pool's 10k tail and is not comparable.